

## Predicting AUS rain occurrence using Logistic Regression

```
df.info()
```

```
from google.colab import drive
```

```
from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
df = pd.read_csv("/content/drive/MyDrive/TIP S.Y.'s/TIP 2021-2022 (3rd  
yr, 1st sem)/Predictive Analytics w Machine Learning/CPE 312 -  
Machine Learning/data/logistic regression/weatherAUS.csv")
```

```
df.head()
```

	Date	Location	MinTemp	...	Temp3pm	RainToday	RainTomorrow
0	2008-12-01	Albury	13.4	...	21.8	No	No
1	2008-12-02	Albury	7.4	...	24.3	No	No
2	2008-12-03	Albury	12.9	...	23.2	No	No
3	2008-12-04	Albury	9.2	...	26.5	No	No
4	2008-12-05	Albury	17.5	...	29.7	No	No

```
[5 rows x 23 columns]
```

### Objective(s):

This activity aims to solve classification problem using logistic regression

```
print(df.shape)
```

```
(145460, 23)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 145460 entries, 0 to 145459  
Data columns (total 23 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Date                  145460 non-null object  
1   Location              145460 non-null object  
2   MinTemp               143975 non-null float64  
3   MaxTemp               144199 non-null float64  
4   Rainfall              142199 non-null float64  
5   Evaporation           82670 non-null float64  
6   Sunshine              75625 non-null float64  
7   WindGustDir           135134 non-null object  
8   WindGustSpeed         135197 non-null float64
```

```

9   WindDir9am      134894 non-null object
10  WindDir3pm      141232 non-null object
11  WindSpeed9am    143693 non-null float64
12  WindSpeed3pm    142398 non-null float64
13  Humidity9am     142806 non-null float64
14  Humidity3pm     140953 non-null float64
15  Pressure9am     130395 non-null float64
16  Pressure3pm     130432 non-null float64
17  Cloud9am        89572 non-null float64
18  Cloud3pm        86102 non-null float64
19  Temp9am         143693 non-null float64
20  Temp3pm         141851 non-null float64
21  RainToday       142199 non-null object
22  RainTomorrow    142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB

```

- Remove null values for RainToday and RainTomorrow since the latter is needed for training and testing and the former feels like a good predictor.

```
df.dropna(subset = ['RainToday', 'RainTomorrow'], inplace=True)
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 140787 entries, 0 to 145458
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Date                140787 non-null object
1   Location            140787 non-null object
2   MinTemp             140319 non-null float64
3   MaxTemp             140480 non-null float64
4   Rainfall            140787 non-null float64
5   Evaporation         81093 non-null float64
6   Sunshine            73982 non-null float64
7   WindGustDir         131624 non-null object
8   WindGustSpeed       131682 non-null float64
9   WindDir9am          131127 non-null object
10  WindDir3pm          137117 non-null object
11  WindSpeed9am        139732 non-null float64
12  WindSpeed3pm        138256 non-null float64
13  Humidity9am         139270 non-null float64
14  Humidity3pm         137286 non-null float64
15  Pressure9am         127044 non-null float64
16  Pressure3pm         127018 non-null float64
17  Cloud9am            88162 non-null float64
18  Cloud3pm            84693 non-null float64
19  Temp9am             140131 non-null float64
20  Temp3pm             138163 non-null float64
21  RainToday           140787 non-null object

```

```
22 RainTomorrow 140787 non-null object
dtypes: float64(16), object(7)
memory usage: 25.8+ MB
```

```
px.histogram(df, x = 'Location', title='Location vs. Rainy Days',
color ="RainToday")
```

Data seems to be evenly distributed. Cities generally get more rains than no. Additionally and generally, most cities experience 20% of the time with rain for the last 10 years.

```
px.histogram(df, x = 'RainTomorrow', title='RainTomorrow vs. Rainy
Days', color ="RainToday")
```

More instances of not raining today and not raining tomorrow than other cases.

So we can observe that AUS predictor for tomorrow not raining is if it did not rain yesterday. There seems to be no correlation to raining today and then raining tomorrow VS NOT raining today and then raining tomorrow. So this means using the RainToday predictor to whether it WILL rain tomorrow would be bad since there is no clear distinction that RainToday values reveals a tendency for it to rain tomorrow.

TLDR: Predictor column 'RainToday' is better at predicting "No" for 'RainTomorrow' than predicting "Yes".

```
px.histogram(df, x = 'Temp3pm', title='Temp at 3pm vs. Rainy Days',
color ="RainToday")
```

We can see that usually there are more instances of "No" for "RainTomorrow" than "Yes" for any given temperature at 3PM. Most of the days experience an average temperature of 15-25 degrees Celsius. It is also around these temperatures that we see the most amount of rainy days, even though it is around these temps that we also see the most amount of non-rainy days. Higher temperatures of 35-45+ almost guarantee no rain.

```
px.scatter(df.sample(5000), x = 'MinTemp', y = 'MaxTemp',
title="Minimum Temperature vs Maximum Temperature as predictor for
tomorrow's rain", color ="RainToday", opacity=0.5)
```

Generally, for the days with similar minimum temperatures, a higher maximum temperature seems to indicate that it will not rain on that day. The vertical shift of the otherwise similarly positioned scatter plot of the two predictor columns (with RainToday as indicated) is due to the fact that the higher maximum temperature is, the more likely to not rain on that day.

## Training, Validation, and Test sets

- training 60%
- validation 20%

- test 20%

```
from sklearn.model_selection import train_test_split

train_val_df, test_df = train_test_split(df, test_size=0.2,
random_state=101)
train_df, val_df = train_test_split(train_val_df, test_size=0.25,
random_state=101)
```

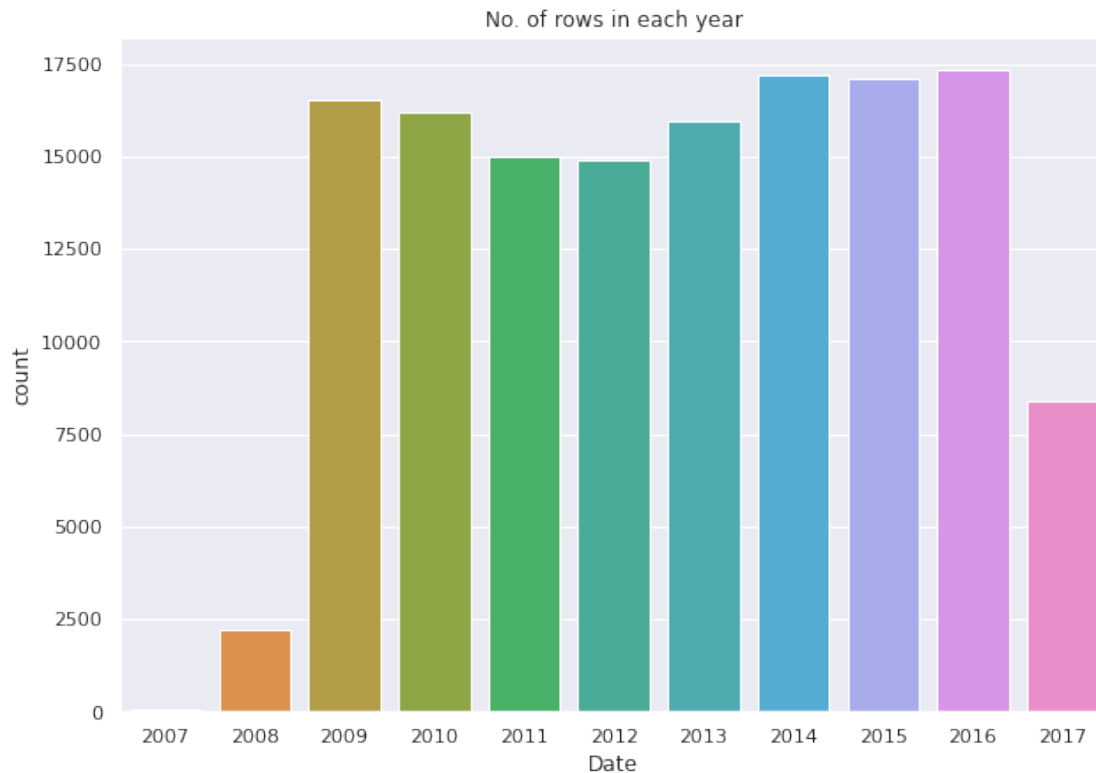
Conclusion:

```
print("shape of train_df is {}".format(train_df.shape))
print("shape of validity_df is {}".format(val_df.shape))
print("shape of test_df is {}".format(test_df.shape))
```

```
shape of train_df is (84471, 23)
shape of validity_df is (28158, 23)
shape of test_df is (28158, 23)
```

Since we are working with data that involves time it is best that we don't train our model with dates that would appear in our validation and test dataset. Meaning one should not train models based on data that is in the future or will be included in the test/validation dataset.

```
%matplotlib inline
fig, ax = plt.subplots(figsize=(10,7))
sns.set_theme(style="darkgrid")
ax = sns.countplot(data = df,x= pd.to_datetime(df["Date"]).dt.year)
ax.set_title("No. of rows in each year");
```



So the data for the test/training/validation will be broken down like this:

1. train\_df will contain data points with years below 2015
2. val\_df will contain data points with years equal to 2015
3. test\_df will contain data points with years greater than 2015

```
year = pd.to_datetime(df["Date"]).dt.year
```

```
train_df = df[year < 2015]
```

```
val_df = df[year == 2015]
```

```
test_df = df[year > 2015]
```

```
print("shape of train_df is {}".format(train_df.shape))
```

```
print("shape of validity_df is {}".format(val_df.shape))
```

```
print("shape of test_df is {}".format(test_df.shape))
```

```
shape of train_df is (97988, 23)
```

```
shape of validity_df is (17089, 23)
```

```
shape of test_df is (25710, 23)
```

```
train_inputs = train_df.drop(['Date', 'RainTomorrow'], axis=1)
```

```
train_target = train_df["RainTomorrow"]
```

```
val_inputs = val_df.drop(['Date', 'RainTomorrow'], axis=1)
```

```
val_target = val_df["RainTomorrow"]
```

```
test_inputs = test_df.drop(['Date', 'RainTomorrow'], axis=1)
```

```
test_target = test_df["RainTomorrow"]
```

collect all the columns that are numerical and categorical and put them into a list

```
num_cols =
train_inputs.select_dtypes(include=np.number).columns.tolist()
cat_cols =
train_inputs.select_dtypes(include=["object"]).columns.tolist()

train_inputs[num_cols].describe()
```

	MinTemp	MaxTemp	...	Temp9am	Temp3pm
count	97674.000000	97801.000000	...	97414.000000	97392.000000
mean	12.007831	23.022202	...	16.835126	21.540138
std	6.347175	6.984397	...	6.404586	6.831612
min	-8.500000	-4.100000	...	-5.900000	-5.100000
25%	7.500000	17.900000	...	12.200000	16.600000
50%	11.800000	22.400000	...	16.600000	20.900000
75%	16.600000	27.900000	...	21.400000	26.200000
max	33.900000	48.100000	...	40.200000	46.100000

[8 rows x 16 columns]

```
train_inputs[cat_cols].nunique()
```

```
Location      49
WindGustDir   16
WindDir9am    16
WindDir3pm    16
RainToday     2
dtype: int64
```

## Imputing missing data

```
from sklearn.impute import SimpleImputer
```

```
#create an imputer object
```

```
imputer = SimpleImputer(strategy="mean")
```

```
# check the columns that have nan values and how many
```

```
train_inputs[num_cols].isnull().sum()
```

```
MinTemp      314
MaxTemp      187
Rainfall      0
Evaporation  36331
Sunshine     40046
WindGustSpeed 6828
WindSpeed9am  874
WindSpeed3pm 1069
Humidity9am   1052
Humidity3pm   1116
Pressure9am   9112
```

```
Pressure3pm      9131
Cloud9am         34988
Cloud3pm         36022
Temp9am          574
Temp3pm          596
dtype: int64
```

```
# fit the imputer model to fill each column with missing values the
mean value for that column
```

```
imputer.fit(train_inputs[num_cols])
```

```
SimpleImputer()
```

```
# the object imputer now contains an attribute called .statistics_
which contains the mean value for each column. We can access this:
```

```
list(imputer.statistics_)
```

```
[12.0078311526097,
 23.0222022269711,
 2.37293546148508,
 5.289991404057933,
 7.609003831417626,
 40.21587318999561,
 14.09226270156723,
 18.76460755888938,
 68.62874473879673,
 51.46954744404988,
 1017.5137337413925,
 1015.132351891241,
 4.302952380952381,
 4.410676822773779,
 16.835126367873205,
 21.54013779365862]
```

```
# we need to inject these values in the predictor variable for all our
datasets.
```

```
train_inputs[num_cols] = imputer.fit_transform(train_inputs[num_cols])
```

```
val_inputs[num_cols] = imputer.fit_transform(val_inputs[num_cols])
```

```
test_inputs[num_cols] = imputer.fit_transform(test_inputs[num_cols])
```

## Normalizing our Numerical columns

```
from sklearn.preprocessing import MinMaxScaler
```

```
#create an object for MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
#scaler.transform({data fram with num cols}) will result in the
scaling of the values from (0,1)
```

```
train_inputs[num_cols] = scaler.fit_transform(train_inputs[num_cols])
```

```

val_inputs[num_cols] = scaler.fit_transform(val_inputs[num_cols])
test_inputs[num_cols] = scaler.fit_transform(test_inputs[num_cols])

# verify that the scaling worked (val_inputs is used to see the max
and min but train_inputs and test_inputs both have a max of 1 and min
of 0)
val_inputs[num_cols].describe().loc[["min", "max"]]

```

	MinTemp	MaxTemp	Rainfall	...	Cloud3pm	Temp9am	Temp3pm
min	0.0	0.0	0.0	...	0.0	0.0	0.0
max	1.0	1.0	1.0	...	1.0	1.0	1.0

[2 rows x 16 columns]

## Encoding Categorical Data

```

# check all of our categorical columns
train_inputs[cat_cols].nunique()

```

```

Location          49
WindGustDir       16
WindDir9am        16
WindDir3pm        16
RainToday         2
dtype: int64

```

```

train_inputs[cat_cols].isnull().sum()

```

```

Location          0
WindGustDir       6868
WindDir9am        7019
WindDir3pm        1952
RainToday         0
dtype: int64

```

```

from sklearn.preprocessing import OneHotEncoder

```

```

# encoder object
encoder = OneHotEncoder(sparse = False, handle_unknown='ignore')

```

## Try to impute missing categorical values using logistic regression

*# Separate the test, train, and val data into two dataframes: one w/ missing data and one w/o.*

## Try to impute missing categorical values by inputting "Unknown"

```

test_inputs.fillna("Unknown", inplace = True)
train_inputs.fillna("Unknown", inplace = True)
val_inputs.fillna("Unknown", inplace = True)

```

*#check to see if imputation works (val\_df2 is shown here but test\_df2 and train\_df2 also are modified)*

```

train_inputs

```



	Location	MinTemp	MaxTemp	...	Temp9am	Temp3pm
RainToday						
0	Albury	0.516509	0.517241	...	0.494577	0.525391
No						
1	Albury	0.375000	0.559387	...	0.501085	0.574219
No						
2	Albury	0.504717	0.570881	...	0.583514	0.552734
No						
3	Albury	0.417453	0.614943	...	0.520607	0.617188
No						
4	Albury	0.613208	0.697318	...	0.514100	0.679688
No						
...	...	...	...	...	...	...
.						
144548	Uluru	0.599057	0.714559	...	0.642082	0.720703
No						
144549	Uluru	0.556604	0.783525	...	0.754881	0.779297
No						
144550	Uluru	0.608491	0.802682	...	0.772234	0.796875
No						
144551	Uluru	0.674528	0.816092	...	0.774403	0.826172
No						
144552	Uluru	0.731132	0.837165	...	0.780911	0.830078
No						

[97988 rows x 21 columns]

*# We must first fit our encoder object with our dataframe so it can recognize the categorical columns*

```
encoder.fit(train_inputs[cat_cols])
```

```
OneHotEncoder(handle_unknown='ignore', sparse=False)
```

```
encoder.categories_
```

```
[array(['Adelaide', 'Albany', 'Albury', 'AliceSprings',
'BadgerysCreek',
'Ballarat', 'Bendigo', 'Brisbane', 'Cairns', 'Canberra',
'Cobar',
'CoffsHarbour', 'Dartmoor', 'Darwin', 'GoldCoast', 'Hobart',
'Katherine', 'Launceston', 'Melbourne', 'MelbourneAirport',
'Mildura', 'Moree', 'MountGambier', 'MountGinini',
'Newcastle',
'Nhil', 'NorahHead', 'NorfolkIsland', 'Nuriootpa',
'PearceRAAF',
'Penrith', 'Perth', 'PerthAirport', 'Portland', 'Richmond',
'Sale',
'SalmonGums', 'Sydney', 'SydneyAirport', 'Townsville',
'Tuggeranong', 'Uluru', 'WaggaWagga', 'Walpole', 'Watsonia',
'Williamtown', 'Witchcliffe', 'Wollongong', 'Woomera'],
dtype=object),
```

```

array(['E', 'ENE', 'ESE', 'N', 'NE', 'NNE', 'NNW', 'NW', 'S', 'SE',
'SSE',
      'SSW', 'SW', 'Unknown', 'W', 'WNW', 'WSW'], dtype=object),
array(['E', 'ENE', 'ESE', 'N', 'NE', 'NNE', 'NNW', 'NW', 'S', 'SE',
'SSE',
      'SSW', 'SW', 'Unknown', 'W', 'WNW', 'WSW'], dtype=object),
array(['E', 'ENE', 'ESE', 'N', 'NE', 'NNE', 'NNW', 'NW', 'S', 'SE',
'SSE',
      'SSW', 'SW', 'Unknown', 'W', 'WNW', 'WSW'], dtype=object),
array(['No', 'Yes'], dtype=object)]

```

```

# generate columns names to label our new one-hot encoded columns
encoded_cols = list(encoder.get_feature_names(cat_cols))

```

```

print(encoded_cols)

```

```

['Location_Adelaide', 'Location_Albury', 'Location_AliceSprings', 'Location_BadgerysCreek',
'Location_Ballarat', 'Location_Bendigo', 'Location_Brisbane',
'Location_Cairns', 'Location_Canberra', 'Location_Cobar',
'Location_CoffsHarbour', 'Location_Dartmoor', 'Location_Darwin',
'Location_GoldCoast', 'Location_Hobart', 'Location_Katherine',
'Location_Launceston', 'Location_Melbourne',
'Location_MelbourneAirport', 'Location_Mildura', 'Location_Moree',
'Location_MountGambier', 'Location_MountGinini', 'Location_Newcastle',
'Location_Nhil', 'Location_NorahHead', 'Location_NorfolkIsland',
'Location_Nuriootpa', 'Location_PearceRAAF', 'Location_Penrith',
'Location_Perth', 'Location_PerthAirport', 'Location_Portland',
'Location_Richmond', 'Location_Sale', 'Location_SalmonGums',
'Location_Sydney', 'Location_SydneyAirport', 'Location_Townsville',
'Location_Tuggeranong', 'Location_Uluru', 'Location_WaggaWagga',
'Location_Walpole', 'Location_Watsonia', 'Location_Williamtown',
'Location_Witchcliffe', 'Location_Wollongong', 'Location_Woomera',
'WindGustDir_E', 'WindGustDir_ENE', 'WindGustDir_ESE',
'WindGustDir_N', 'WindGustDir_NE', 'WindGustDir_NNE',
'WindGustDir_NNW', 'WindGustDir_NW', 'WindGustDir_S',
'WindGustDir_SE', 'WindGustDir_SSE', 'WindGustDir_SSW',
'WindGustDir_SW', 'WindGustDir_Unknown', 'WindGustDir_W',
'WindGustDir_WNW', 'WindGustDir_WSW', 'WindDir9am_E',
'WindDir9am_ENE', 'WindDir9am_ESE', 'WindDir9am_N', 'WindDir9am_NE',
'WindDir9am_NNE', 'WindDir9am_NNW', 'WindDir9am_NW', 'WindDir9am_S',
'WindDir9am_SE', 'WindDir9am_SSE', 'WindDir9am_SSW', 'WindDir9am_SW',
'WindDir9am_Unknown', 'WindDir9am_W', 'WindDir9am_WNW',
'WindDir9am_WSW', 'WindDir3pm_E', 'WindDir3pm_ENE', 'WindDir3pm_ESE',
'WindDir3pm_N', 'WindDir3pm_NE', 'WindDir3pm_NNE', 'WindDir3pm_NNW',
'WindDir3pm_NW', 'WindDir3pm_S', 'WindDir3pm_SE', 'WindDir3pm_SSE',
'WindDir3pm_SSW', 'WindDir3pm_SW', 'WindDir3pm_Unknown',
'WindDir3pm_W', 'WindDir3pm_WNW', 'WindDir3pm_WSW', 'RainToday_No',
'RainToday_Yes']

```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/  
deprecation.py:87: FutureWarning:
```

Function `get_feature_names` is deprecated; `get_feature_names` is deprecated in 1.0 and will be removed in 1.2. Please use `get_feature_names_out` instead.

```
# Now that we have generated the one-hot encoded columns, we shall  
append them to our train, validation, and test datasets  
train_inputs[encoded_cols] = encoder.fit_transform(train_df[cat_cols])  
test_inputs[encoded_cols] = encoder.fit_transform(test_df[cat_cols])  
val_inputs[encoded_cols] = encoder.fit_transform(val_df[cat_cols])
```

```
# Delete redundant columns such as the categorical columns such as the  
ones we used the one-encoded on  
train_inputs.drop(cat_cols, axis = 1, inplace = True)  
test_inputs.drop(cat_cols, axis = 1, inplace = True)  
val_inputs.drop(cat_cols, axis = 1, inplace = True)
```

```
#same outputs for the train and validation datasets.  
test_inputs.columns.tolist()
```

```
['MinTemp',  
 'MaxTemp',  
 'Rainfall',  
 'Evaporation',  
 'Sunshine',  
 'WindGustSpeed',  
 'WindSpeed9am',  
 'WindSpeed3pm',  
 'Humidity9am',  
 'Humidity3pm',  
 'Pressure9am',  
 'Pressure3pm',  
 'Cloud9am',  
 'Cloud3pm',  
 'Temp9am',  
 'Temp3pm',  
 'Location_Adelaide',  
 'Location_Albury',  
 'Location_Albury',  
 'Location_AliceSprings',  
 'Location_BadgerysCreek',  
 'Location_Ballarat',  
 'Location_Bendigo',  
 'Location_Brisbane',  
 'Location_Cairns',  
 'Location_Canberra',  
 'Location_Cobar',  
 'Location_CoffsHarbour',
```

'Location\_Dartmoor',  
'Location\_Darwin',  
'Location\_GoldCoast',  
'Location\_Hobart',  
'Location\_Katherine',  
'Location\_Launceston',  
'Location\_Melbourne',  
'Location\_MelbourneAirport',  
'Location\_Mildura',  
'Location\_Moree',  
'Location\_MountGambier',  
'Location\_MountGinini',  
'Location\_Newcastle',  
'Location\_Nhil',  
'Location\_NorahHead',  
'Location\_NorfolkIsland',  
'Location\_Nuriootpa',  
'Location\_PearceRAAF',  
'Location\_Penrith',  
'Location\_Perth',  
'Location\_PerthAirport',  
'Location\_Portland',  
'Location\_Richmond',  
'Location\_Sale',  
'Location\_SalmonGums',  
'Location\_Sydney',  
'Location\_SydneyAirport',  
'Location\_Townsville',  
'Location\_Tuggeranong',  
'Location\_Uluru',  
'Location\_WaggaWagga',  
'Location\_Walpole',  
'Location\_Watsonia',  
'Location\_Williamtown',  
'Location\_Witchcliffe',  
'Location\_Wollongong',  
'Location\_Woomera',  
'WindGustDir\_E',  
'WindGustDir\_ENE',  
'WindGustDir\_ESE',  
'WindGustDir\_N',  
'WindGustDir\_NE',  
'WindGustDir\_NNE',  
'WindGustDir\_NNW',  
'WindGustDir\_NW',  
'WindGustDir\_S',  
'WindGustDir\_SE',  
'WindGustDir\_SSE',  
'WindGustDir\_SSW',  
'WindGustDir\_SW',

```
'WindGustDir_Unknown',
'WindGustDir_W',
'WindGustDir_WNW',
'WindGustDir_WSW',
'WindDir9am_E',
'WindDir9am_ENE',
'WindDir9am_ESE',
'WindDir9am_N',
'WindDir9am_NE',
'WindDir9am_NNE',
'WindDir9am_NNW',
'WindDir9am_NW',
'WindDir9am_S',
'WindDir9am_SE',
'WindDir9am_SSE',
'WindDir9am_SSW',
'WindDir9am_SW',
'WindDir9am_Unknown',
'WindDir9am_W',
'WindDir9am_WNW',
'WindDir9am_WSW',
'WindDir3pm_E',
'WindDir3pm_ENE',
'WindDir3pm_ESE',
'WindDir3pm_N',
'WindDir3pm_NE',
'WindDir3pm_NNE',
'WindDir3pm_NNW',
'WindDir3pm_NW',
'WindDir3pm_S',
'WindDir3pm_SE',
'WindDir3pm_SSE',
'WindDir3pm_SSW',
'WindDir3pm_SW',
'WindDir3pm_Unknown',
'WindDir3pm_W',
'WindDir3pm_WNW',
'WindDir3pm_WSW',
'RainToday_No',
'RainToday_Yes']
```

```
print("train_inputs shape is {}".format(train_inputs.shape))
print("train_target shape is {}".format(train_target.shape))
print("val_inputs shape is {}".format(val_inputs.shape))
print("val_target shape is {}".format(val_target.shape))
print("test_inputs shape is {}".format(test_inputs.shape))
print("test_target shape is {}".format(test_target.shape))
```

```
train_inputs shape is (97988, 118)
train_target shape is (97988,)
val_inputs shape is (17089, 118)
```

```
val_target shape is (17089,)
test_inputs shape is (25710, 118)
test_target shape is (25710,)
```

### Train the training set using Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression(solver='liblinear')
```

```
model.fit(train_inputs, train_target)
```

```
LogisticRegression(solver='liblinear')
```

```
weight_df = pd.DataFrame(
    {"feature": train_inputs.columns.tolist(),
     "weight": model.coef_.tolist()[0]}
)
weight_df
```

	feature	weight
0	MinTemp	0.894303
1	MaxTemp	-2.854273
2	Rainfall	3.162893
3	Evaporation	0.640867
4	Sunshine	-1.649886
...	...	...
113	WindDir3pm_W	-0.043347
114	WindDir3pm_WNW	-0.285430
115	WindDir3pm_WSW	0.068019
116	RainToday_No	-1.559546
117	RainToday_Yes	-1.060584

```
[118 rows x 2 columns]
```

```
import plotly.graph_objects as go
```

```
neg_weights = weight_df[weight_df["weight"] < 0].sort_values("weight",
    ascending = False)
pos_weights = weight_df[weight_df["weight"] > 0].sort_values("weight",
    ascending = False)
```

```
fig = go.Figure()
fig.add_trace(go.Bar(x=neg_weights.feature, y=neg_weights.weight,
    marker_color='crimson',
    name='negative weights'))
fig.add_trace(go.Bar(x=pos_weights.feature, y=pos_weights.weight,
    marker_color='lightslategrey',
    name='positive weights'
    ))
```

```
#fig.update_xaxes(visible = False)
```

```
fig.update_xaxes(nticks = 100)
```

```
fig.update_yaxes(nticks=20, ticks = "outside")
fig.update_layout(
    margin=dict(l=20, r=20, t=20, b=20),
    paper_bgcolor="LightSteelBlue"
)
fig.show()

'''
press on the legends to see/remove the negative and positive weights.
Negative Weights is turned off currently. press the negative weights
legend to show the negative weights
'''
```

```
{"type": "string"}
```

generally the higher the coefficient the more the feature is able to accurately predict the target variable. A high negative coefficient value means that the relationship is inverse, but strong.

### Making Predictions on the validation set

```
val_predictions = model.predict(val_inputs)
val_predictions.tolist()
```

```
# check the accuracy of the train_target and train_predictions
from sklearn.metrics import accuracy_score
print("accuracy is: {}".format(accuracy_score(val_target,
val_predictions)))
```

```
accuracy is: 0.8516004447305284
```

```
# check the prediction probabilities for the observations
val_proba = model.predict_proba(val_inputs)
print(model.classes_)
print(val_proba)
```

```
['No' 'Yes']
[[0.98930488 0.01069512]
 [0.96371828 0.03628172]
 [0.94248961 0.05751039]
 ...
 [0.97879253 0.02120747]
 [0.9711122  0.0288878 ]
 [0.94725931 0.05274069]]
```

```
from sklearn.metrics import confusion_matrix
confusion_matrix(val_target, val_predictions, normalize = 'true')
```


```
array([[0.96661979, 0.03338021],
       [0.58272778, 0.41727222]])
```

```
data = confusion_matrix(val_target, val_predictions, normalize =
'true')
```

```

df_cm = pd.DataFrame(data, columns=np.unique(val_predictions), index =
np.unique(val_target))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
plt.figure(figsize = (5,5))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, cmap="Blues", annot=True,annot_kws={"size": 16});

```



as one can see the confusion matrix reveals that the model is better at predicting whether it will not rain tomorrow than it can predict whether it will rain tomorrow.

```

from sklearn.metrics import classification_report
print(classification_report(val_target, val_predictions))

```

	precision	recall	f1-score	support
No	0.86	0.97	0.91	13511
Yes	0.77	0.42	0.54	3578
accuracy			0.85	17089
macro avg	0.82	0.69	0.73	17089



weighted avg	0.84	0.85	0.83	17089
--------------	------	------	------	-------

### first attempt at optimizing model by changing the hyperparameters of the logistic regression model

```
param_grid = [
    {'penalty': ['l1', 'l2', 'elasticnet', 'none'],
     'solver': ['liblinear', 'newton-cg', 'lbfgs', 'sag', 'saga'],
     'max_iter': [100, 1000, 1500, 2000]}
]

from sklearn.model_selection import GridSearchCV

logmodel = LogisticRegression()

clf = GridSearchCV(logmodel, param_grid, cv=2, verbose=True, n_jobs=-1)

best_clf = clf.fit(train_inputs, train_target)
```

Fitting 2 folds for each of 80 candidates, totalling 160 fits

```
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:
```

72 fits failed out of a total of 160.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting `error_score='raise'`.

Below are more details about the failures:

-----  
-----

8 fits failed with the following error:

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 1461, in fit
```

```
    solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 449, in _check_solver
```

```
    % (solver, penalty)
```

```
ValueError: Solver newton-cg supports only 'l2' or 'none' penalties,
```

got l1 penalty.

```
-----  
-----  
8 fits failed with the following error:  
Traceback (most recent call last):  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 1461, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 449, in _check_solver  
    % (solver, penalty)  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
-----  
-----  
8 fits failed with the following error:  
Traceback (most recent call last):  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 1461, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 449, in _check_solver  
    % (solver, penalty)  
ValueError: Solver sag supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
-----  
-----  
8 fits failed with the following error:  
Traceback (most recent call last):  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File  
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic
```

```
.py", line 1461, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic
.py", line 459, in _check_solver
    solver
ValueError: Only 'saga' solver supports elasticnet penalty, got
solver=liblinear.
```

```
-----
-----
8 fits failed with the following error:
Traceback (most recent call last):
File
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_valid
ation.py", line 681, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic
.py", line 1461, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic
.py", line 449, in _check_solver
    % (solver, penalty)
ValueError: Solver newton-cg supports only 'l2' or 'none' penalties,
got elasticnet penalty.
```

```
-----
-----
8 fits failed with the following error:
Traceback (most recent call last):
File
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_valid
ation.py", line 681, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic
.py", line 1461, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic
.py", line 449, in _check_solver
    % (solver, penalty)
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got
elasticnet penalty.
```

```
-----
-----
8 fits failed with the following error:
```

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 1461, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 449, in _check_solver
    % (solver, penalty)
ValueError: Solver sag supports only 'l2' or 'none' penalties, got elasticnet penalty.
```

-----  
-----

8 fits failed with the following error:

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 1473, in fit
    % self.l1_ratio
ValueError: l1_ratio must be between 0 and 1; got (l1_ratio=None)
```

-----  
-----

8 fits failed with the following error:

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 681, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 1461, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File
"/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py", line 464, in _check_solver
    raise ValueError("penalty='none' is not supported for the liblinear solver")
ValueError: penalty='none' is not supported for the liblinear solver
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_search.py:972: UserWarning:
```

```
One or more of the test scores are non-finite: [0.821968          nan
nan          nan 0.82289668 0.82062089
 0.822366    0.82195779 0.8223558    0.82234559          nan          nan
          nan          nan          nan          nan 0.78861697 0.81428338
0.80577214 0.81313018 0.82187615          nan          nan          nan
0.82289668 0.82062089 0.822366    0.82232518 0.822366    0.82234559
          nan          nan          nan          nan          nan          nan
0.78861697 0.80469037 0.79499531 0.80854799 0.82166184          nan
          nan          nan 0.82289668 0.82062089 0.822366    0.82232518
0.8223558    0.82233539          nan          nan          nan          nan
          nan          nan 0.78861697 0.80469037 0.80569049 0.80853778
0.82174348          nan          nan          nan 0.82289668 0.82062089
0.822366    0.82232518 0.82237621 0.82234559          nan          nan
          nan          nan          nan          nan 0.78861697 0.80469037
0.80542515 0.80854799]
```

```
best_clf.best_params_
```

```
{'max_iter': 100, 'penalty': 'l1', 'solver': 'saga'}
```

```
print("previous accuracy is: {}".format(accuracy_score(val_target,
val_predictions)))
logmodel_optimized = LogisticRegression(max_iter = 100,penalty='l1',
solver='saga').fit(train_inputs, train_target)
```

```
previous accuracy is: 0.8516004447305284
```

```
val_predictions_optimized = logmodel_optimized.predict(val_inputs)
print("new accuracy is: {}".format(accuracy_score(val_target,
val_predictions_optimized)))
```

```
new accuracy is: 0.85171747908011
```

```
print("difference in accruacy scores are:
{}".format(accuracy_score(val_target, val_predictions_optimized) -
accuracy_score(val_target, val_predictions)))
```

```
difference in accruacy scores are: 0.00011703434958154624
```

The logistical regression model was optimized by changing the hyperparameters. However, after several fits using the hyperparameter combinations of param\_grid, the accuracy score did not change that much. There is an improvement of 0.00011703434958154624.


Now we wil check the confusion matrix

```
data = confusion_matrix(val_target, val_predictions_optimized,
normalize = 'true')
df_cm = pd.DataFrame(data,
```

```

columns=np.unique(val_predictions_optimized), index =
np.unique(val_target))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
plt.figure(figsize = (5,5))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, cmap="Blues", annot=True,annot_kws={"size": 16});

```



looks like the model's ability to predict certainty of actually raining decreased.

**Evaluate the FINAL model (using test dataset) using classification report, accuracy and confusion matrix**

```

test_predictions = logmodel_optimized.predict(test_inputs)
print(classification_report(test_target,test_predictions))

```

	precision	recall	f1-score	support
No	0.85	0.96	0.90	19885
Yes	0.77	0.43	0.55	5825
accuracy			0.84	25710
macro avg	0.81	0.70	0.73	25710

weighted avg          0.83          0.84          0.82          25710

```
print(accuracy_score(test_target, test_predictions))
```

```
0.841267989109296
```

```
print(confusion_matrix(test_target, test_predictions, normalize =  
'true'))
```

```
[[0.96208197 0.03791803]  
 [0.5711588  0.4288412 ]]
```

```
data = confusion_matrix(test_target, test_predictions, normalize =  
'true')
```

```
df_cm = pd.DataFrame(data, columns=np.unique(test_predictions), index  
= np.unique(test_target))
```


```
df_cm.index.name = 'Actual'
```

```
df_cm.columns.name = 'Predicted'
```

```
plt.figure(figsize = (5,5))
```

```
sns.set(font_scale=1.4)#for label size
```

```
sns.heatmap(df_cm, cmap="Blues", annot=True,annot_kws={"size": 16});
```



## type your conclusion here

In conclusion the resulting accuracy score, confusion matrix, and classification report are undescribable in my opinion. I have initially tried to optimize the model by increasing the combinations of the hyperparameters; however, the processing of determining which combination was the best took so much time. When the process was not done until 3 hours, I repeated the entire process again using lesser parameters. I went from 4800 fits to 160 fits using GridSearchCV.

According to the confusion matrix, the model is better at predicting No rain than it is at predicting that it would rain. In fact, if the model were used in weather forecasts and it outputs a value of "Yes", people are slightly likely to see it not rain as opposed to not raining. People should be more confident that it would not rain if the model tells them it would not rain.